

Edwin Zhang

Researcher
OpenAI

Email: edwin@openai.com
Web URL: <https://eddie.win/>
Google Scholar: <https://scholar.google...>

Education

- **Harvard University**
Ph.D., Computer Science, RL for Social Good, 2023 - (indefinite leave)
Advisor: Milind Tambe
- **University of California, Santa Barbara**
M.S., Computer Science, 2022 - 2023
Advisor: William Wang
- **University of California, Santa Barbara**
B.S., Computer Science, 2019 - 2022
Graduated top of class at age 19
Advisor: William Wang
High Honors, Regents Scholar (Top 2.5% of school)
- **Westview High School**
Graduated top of class at age 16, 2018 - 2019

Research Interests

AI for Policy; Real-world AI deployment; Economic impacts of AI; AI Safety and Alignment; Principles of intelligence; Reinforcement learning

Experience

- **OpenAI**, CA 07/2024 – Present.
Member of Technical Staff
 - RL Post-Training, Collective Alignment, and Hallucinations.
- **Humanity Unleashed**, MA 09/2023 – Present.
Technical Director
 - Working on something exciting.
- **Harvard University** (on leave), MA 09/2023 – 07/2024.
Graduate Student Researcher
 - Creating new foundation models for social good through proposing and implementing a new Language-Conditioned Restless Multi-Armed Bandits (RMABs) framework.
 - Working on improving world's largest maternal health care program, positively affecting over 10 million mothers in India (Kilkari and mMitra).

- **Massachusetts Institute of Technology-IBM Watson AI**, Cambridge, MA 12/2022 – 05/2023.
Visiting Researcher, Mentors: Chuang Gan and Yikang Shen
 - Started work on an ongoing hierarchical RL project with Chuang Gan with potential for solving extremely long-horizon control problems such as minecraft diamond crafting.
 - Won 3rd place out of 19 in NeurIPS Integrated Language and Understanding Challenge, receiving a \$1500 cash prize.
- **Meta**, Seattle, WA 07/2022 – 09/2022.
Research Intern, Mentor: Amy Zhang
 - Proposed, analyzed, and deployed new group page configuration reducing misinformation by 4%, improving the experience of 3 million daily active users
 - Created new facebook post ranking model with 17% gain on offline engagement area under curve (AUC) metrics
 - Started work on Language Control Diffusion.
- **Plato Systems**, Palo Alto, CA 06/2021 – 06/2022.
Computer Vision and Software Engineering Intern, Mentor: Mashour Solh
 - Developed multiple view calibration pipeline through planar homographies and OpenCV.
 - Created set up process and capture script for NVIDIA Jetson platform with multiple third party imaging providers.
- **Allthenticate**, Santa Barbara, CA 06/2020 – 06/2021.
Lead Full Stack Engineer/First Hire, Mentor: Chad Spensky
 - Led development on cloud platform in early stage startup, collaborating directly with the CEO to architect and implement proprietary API.
 - Taught advanced Vue JS by taking complete responsibility at each step of the development phase – delivered a full web application while teaching and leading two other interns working on the same project.
 - Built and deployed python backend with over 27000 lines to use Elastic Beanstalk, implementing dockerized development process to speed up iteration cycles by 25%.
- **Yaitea**, CA 08/2018 – 09/2019.
Founder and Lead Tutor
 - Assessed a need for tutoring code and critical thinking to children, as programming skills arose in demand and traditional tutoring services struggled to keep up.
 - Collaborated with several students and parents to create lasting relationships.

Publications and Preprints

In Submission

1. *Edward S. Kang, *Hoon Shin, *Felix Chen, *Benjamin Smith, *Shrenik Patel, Xingzhe Li, Defu Cao, Christoph Bergmeir, **Edwin Zhang**. "ORBITS: Online Real-time Benchmarking of Indicator Time Series." *In submission to Neural Information Processing Systems, 2025.*

Preprints

1. OpenAI. "OpenAI o1 System Card." *arXiv*, 2024.

Published

1. Natalie Abreu, **Edwin Zhang**, Eran Malach, Naomi Saphra. "A Taxonomy of Transcendence." *Second Conference on Language Modeling*, 2025.
2. Matthew Ho, Vincent Zhu, Xiaoyin Chen, Moksh Jain, Nikolay Malkin, **Edwin Zhang**. "Proof Flow: Preliminary Study on Generative Flow Network Language Model Tuning for Formal Reasoning." *First Workshop on System-2 Reasoning at Scale, Neural Information Processing Systems*, 2024
3. **Edwin Zhang**, Naomi Saphra, Vincent Zhu, Anat Kleiman, Benjamin L. Edelman, Milind Tambe, Sham M. Kakade, Eran Malach. "Transcendence: Generative Models Can Outperform The Experts That Train Them." *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
4. *Nikhil Behari, ***Edwin Zhang**, Yunfan Zhao, Aparna Taneja, Dheeraj Nagaraj, Milind Tambe. "A Decision-Language Model (DLM) for Dynamic Restless Multi-Armed Bandit Tasks in Public Health." *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
5. *Danqing Wang, *Antonis Antoniadis, Kha-Dinh Luong, **Edwin Zhang**, Mert Kosan, Jiachen Li, Ambuj Singh, William Yang Wang, Lei Li. "Global Human-guided Counterfactual Explanations for Molecular Properties via Reinforcement Learning." *International Conference on Knowledge Discovery and Data Mining*, 2024.
6. **Edwin Zhang**, Sadie Zhao, Tonghan Wang, Safwan Hossain, Stephan Zheng, David C. Parkes, Milind Tambe, and Yiling Chen. "Social Environment Design." *International Conference on Machine Learning 2024, Position Paper*.
7. *Yunfan Zhao, *Nikhil Behari, Edward Hughes, **Edwin Zhang**, Dheeraj Nagaraj, Karl Tuyls, Aparna Taneja and Milind Tambe. "Towards a Pretrained Model for Restless Bandits via Multi-arm Generalization." *International Joint Conference on Artificial Intelligence 2024*.
8. *Lauren Cooke, *Harvey Klyne, ***Edwin Zhang**, Cassidy Laidlaw, Milind Tambe, and Finale Doshi-Velez. "Toward Computationally Efficient Inverse Reinforcement Learning via Reward Shaping." *International Conference on Learning Representations 2024, Tiny Papers*.
9. **Edwin Zhang**, Yujie Lu, William Wang, and Amy Zhang. "Language Control Diffusion: Efficiently Scaling through Space, Time, and Tasks." *International Conference on Learning Representations 2024*.
10. *Yunfan Zhao, *Nikhil Behari, Edward Hughes, **Edwin Zhang**, Dheeraj Nagaraj, Karl Tuyls, Aparna Taneja and Milind Tambe. "Towards Zero Shot Learning in Restless Multi-armed Bandits." *International Conference on Autonomous Agents and Multiagent Systems 2024 Extended Abstract*.
11. **Edwin Zhang**. "Towards Generalist Agents through Scaling Offline Reinforcement Learning." *Proquest, University of California*, 29 Aug. 2023, escholarship.org/uc/item/69j3d2w2.
12. *Li, Jiachen, ***Edwin Zhang**, Ming Yin, Qinxun Bai, Yu-Xiang Wang, and William Yang Wang. "Offline reinforcement learning with closed-form policy improvement operators." *International Conference on Machine Learning 2023*, pp. 20485-20528. PMLR, 2023.
13. Kiseleva, Julia, Alexey Skrynnik, Artem Zholus, Shrestha Mohanty, Negar Arabzadeh, Marc-Alexandre Côté, ... **Edwin Zhang** ... Mohammad Aliannejadi et al. "Interactive Grounded Language Understanding in a Collaborative Environment: Retrospective on Iglu 2022 Competition." In *NeurIPS 2022 Competition Track*, pp. 204-216. PMLR, 2022.

14. **Edwin Zhang**, Yujie Lu, William Wang, and Amy Zhang. "LAD: Language Augmented Diffusion for Reinforcement Learning." In *Neural Information Processing Systems 2022, Language and Reinforcement Learning Workshop*.

Awards and Honors

- Eric and Susan Dunn Innovation Fund for Engineering and Applied Sciences, 2024
- Google Cloud Research Grant (\$87,200), 2023
- Harvard AI Safety Technical Fellowship, 2023
- Harvard Effective Altruism Precipice Fellowship, 2023
- Second out of 10 in Amazon Alexa Simbot Challenge (\$100,000 cash prize), 2023
- Third out of 19 in Integrated Language and Understanding (IGLU) Challenge at NeurIPS (\$1500 cash prize), 2022
- First out of 16 in React Category at SBhacks, 2022
- UCSB Distinction in the Major: Research Track, 2022
- UCSB High Honors (Top 8.5% of school), 2022
- First out of 78 in Startup Category at SDhacks, 2021
- Best use of Google Cloud out of 71 at SBhacks, 2021
- First overall out of 6 at Santa Barbara Startup Weekend, 2020
- First out of 70 in Database Category at SBhacks, 2020
- Second out of 85 in AI classification competition at UCSB, 2020
- Google Cloud Cybersecurity Grant Winner (\$1000), 2020
- Regents Scholar UCSB (\$24000), 2019
- Google Cloud Startup Grant Winner (\$3000), 2019
- AP Scholar with Distinction, 2019

Invited Talks

1. International Conference of Machine Learning, 7/16/26, Societal Implications of AGI.
2. Stanford University AI Safety Seminar, 4/3/25, Deliberative Alignment: On Alignment through Reasoning in LLMs.
3. Harvard SEAS AI for Social Good Seminar, 10/26/23. Language Control Diffusion: Efficiently Scaling to Generalist Agents through Space, Time, and Tasks.
4. Harvard Law School Tax Law Seminar, 10/25/23. Review and Analysis of the AI Economist: Improving Equality and Productivity with AI-Driven Tax Policies.

5. UCSB Master's Thesis Defense, 5/31/23. Towards Generalist Agents through Scaling Offline Reinforcement Learning. Slides. Video.
6. MIT Brain and Cognitive Sciences, 12/13/22. Integrating Language into Reinforcement Learning through Diffusion.
7. NeurIPS Integrated Grounded Language and Understanding Competition, 12/06/22. Hierarchical RL through Diffusion Models.
8. UCSB Computer Science Research Poster Session, 6/03/22. Offline RL with CFPI.
9. UCSB Natural Language Processing Lab, 4/08/22. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor Review.

Student Mentoring

Parentheses denote first position after mentorship.

- Henry Gasztowtt (Isara CEO, 2025)
- Vincent Zhu (TikTok MLE, 2025)
- Wesley Kam (UCSB BS, 2024)
- Peiyang Song (Caltech BS, 2024)
- Matthew Ho (UCSD PhD, 2024)
- Lauren Cooke (Harvard BS, 2024)
- Gabriel Totev (UCSB BS, 2024)
- Shinda Huang (Amazon SWE, 2023)
- Katelyn Zhang (Google SWE, 2022)
- Yuhao Zhang (Amazon SWE, 2021)

Teaching Experience

- UC Santa Barbara, Spring 2023
Teaching Assistant, *Machine Learning: Upper-division*, 80 students.
- UC Santa Barbara, Winter 2020
Learning Assistant, *Introduction to Computer Science: Lower-division*, 201 students.

Service

- NeurIPS 2024 Sys-2 Workshop Reviewer, 2024
- AAMAS Reviewer, 2023